

An Adaptive Funnel-Based Trading Rule with Transaction-Cost-Safe Exits

Louis Charistias Saragih

May 2026

Abstract

This paper formulates a discrete-time, long-only trading rule based on a drift-adjusted predictive funnel for log prices. The strategy enters and exits positions using standardized price deviations, local momentum, adaptive drift and volatility estimates, and an explicit transaction-cost profit threshold. The main formal property is a no-realized-loss guarantee for completed voluntary exits: whenever the algorithm sells under its own exit rule, the sale price is high enough to cover proportional buy and sell costs. This guarantee is deliberately narrow. It does not remove mark-to-market drawdown, opportunity cost, indefinite holding risk, model error, or finite-horizon liquidation risk. The implementation also includes practical modifications used in the simulation: bounded diagnostic funnels, trailing-profit exits, trend-continuation re-entry, walk-forward parameter selection, online parameter smoothing, and real market bar data with horizon-dependent sampling intervals.

1 Motivation

The project studies a rule-based version of a common trading idea: hold an asset until price has moved favorably, avoid selling at a realized loss, and re-enter only when the market appears statistically attractive or trend participation is justified. The objective is not to prove global optimality or profitability. The objective is to define a mathematically consistent algorithm, identify exactly what is guaranteed, and evaluate the remaining risks using simulation and real market data.

The strategy is built from four ingredients:

1. a predictive funnel for log prices;
2. a momentum filter for timing entries and exits;
3. a transaction-cost-adjusted profit threshold for voluntary sells;
4. adaptive estimation and walk-forward testing to reduce lookahead bias.

The central distinction is between realized wealth and mark-to-market wealth. The strategy can enforce a no-realized-loss condition on completed voluntary exits, but it cannot prevent the open position from losing value before an exit occurs.

2 Market Model and Predictive Funnel

Let $S_t > 0$ be the asset price observed at discrete times $t = 0, 1, \dots, T$, and let

$$X_t = \log S_t$$

be the log price. The baseline simulation model is

$$X_{t+1} = X_t + \mu + \sigma \varepsilon_{t+1}, \quad \varepsilon_{t+1} \sim \mathcal{N}(0, 1),$$

where $\mu \in \mathbb{R}$ and $\sigma > 0$. In the implemented strategy, the trader does not need to know the true μ and σ . Instead, the funnel uses one-step-ahead estimates $\hat{\mu}_{s|s-1}$ and $\hat{\sigma}_{s|s-1}^2$.

After a reference time τ , define

$$C_\tau(t) = X_\tau + \sum_{s=\tau+1}^t \hat{\mu}_{s|s-1}, \quad V_\tau(t) = \sum_{s=\tau+1}^t \hat{\sigma}_{s|s-1}^2, \quad t > \tau.$$

The standardized funnel deviation is

$$Z_{\tau,t} = \frac{X_t - C_\tau(t)}{\sqrt{V_\tau(t)}}.$$

For $k > 0$, the upper and lower log-funnel boundaries are

$$U_\tau(t) = C_\tau(t) + k\sqrt{V_\tau(t)}, \quad L_\tau(t) = C_\tau(t) - k\sqrt{V_\tau(t)}.$$

Thus $Z_{\tau,t} > k$ is equivalent to $X_t > U_\tau(t)$, and $Z_{\tau,t} < -k$ is equivalent to $X_t < L_\tau(t)$.

Remark 1. *At a fixed time under the correctly specified Gaussian baseline, $Z_{\tau,t}$ has a standard normal interpretation. The trading rule, however, monitors the statistic sequentially. Hence k is best viewed as a boundary-width tuning parameter, not as a simple fixed-time significance level.*

3 Momentum and Transaction Costs

For an integer window $\delta \geq 1$, define the finite-difference momentum statistic

$$M_t = X_t - X_{t-\delta}.$$

Positive M_t indicates recent upward movement, while $M_t \leq 0$ indicates that recent upward momentum has stalled or reversed. More precisely, $M_t \leq 0$ means that the log price has not increased relative to $t - \delta$; it does not imply monotone decline during every intermediate step. A generalized momentum version is discussed in the Extensions section.

Suppose buying incurs proportional cost $c_{\text{buy}} \geq 0$ and selling incurs proportional cost $c_{\text{sell}} \in [0, 1)$. A trade entered at τ_b and exited at t has positive after-cost realized multiplier exactly when

$$\frac{(1 - c_{\text{sell}})S_t}{(1 + c_{\text{buy}})S_{\tau_b}} > 1.$$

Equivalently,

$$X_t - X_{\tau_b} > h, \quad h = \log \left(\frac{1 + c_{\text{buy}}}{1 - c_{\text{sell}}} \right).$$

The quantity h is the minimum log return required to break even after costs.

4 Trading Rule

The strategy is long-only. In the visualization, the simulation starts with an initial position at the first tradable observation so that the exit logic can be seen immediately. After a sale, the rule alternates between cash and long exposure.

4.1 Voluntary Exit

Let τ_b be the current entry time. The mathematical upper-funnel exit rule is

$$\tau_s = \inf \{t > \tau_b + \delta : Z_{\tau_b, t} > k, \quad M_t \leq 0, \quad X_t - X_{\tau_b} > h\}.$$

The first condition says price is high relative to the predictive funnel; the second says recent upward momentum has weakened; the third ensures the trade is profitable after transaction costs.

The Extensions section describes an optional generalized momentum variant that modifies only the second condition while leaving the funnel and profit-buffer conditions unchanged.

4.2 Bounded Diagnostic Funnel

In long trends, the cumulative variance $V_{\tau_b}(t)$ can grow so large that the funnel becomes too wide to trigger practical decisions. The implementation therefore uses a bounded diagnostic reference for signal generation. Given a maximum lookback L , define

$$\rho_t = \max\{\tau_b, t - L\}.$$

The diagnostic statistic used for the funnel test is

$$Z_t^{(L)} = \frac{X_t - C_{\rho_t}(t)}{\sqrt{V_{\rho_t}(t)}}.$$

The original entry price X_{τ_b} is still used for profit accounting. In the current implementation, the default is $L = 60$ bars.

4.3 Trailing-Profit Exit

The implementation also allows a trailing-profit exit. Let

$$H_t = \max_{\tau_b \leq u \leq t} X_u$$

be the highest log price since entry. For a trailing threshold $a > 0$, a trailing exit is allowed when

$$H_t - X_t \geq a \quad \text{and} \quad X_t - X_{\tau_b} > h.$$

The default implementation uses $a = 0.04$ in log-price units. Since the profit condition is retained, this practical exit does not weaken the no-realized-loss theorem.

4.4 Re-Entry

After a sale at τ_s , the strategy holds cash and uses τ_s as the new reference point. The original symmetric re-entry rule buys after a lower-funnel recovery:

$$Z_{\tau_s, t} < -k, \quad M_t \geq 0.$$

This is a bargain-entry rule. It works naturally in mean-reverting markets but can remain in cash during persistent upward trends. The implementation therefore adds a trend-continuation entry. For $z_{\text{trend}} \geq 0$, buy when either

$$\tau_b^{\text{next}} = \inf \{t > \tau_s + \delta : (Z_{\tau_s, t} < -k, \quad M_t \geq 0) \quad \text{or} \quad (Z_{\tau_s, t} > z_{\text{trend}}, \quad M_t > 0)\}.$$

The default value is $z_{\text{trend}} = 0.35$, intentionally smaller than k . The first condition seeks a recovered discount; the second permits participation in an ongoing upward trend. At the current Layer 2 stage, z_{trend} is treated as a fixed design constant; in later calibration work it should either be selected by walk-forward validation or made time-varying with the entry/exit thresholds.

5 No-Realized-Loss Property

Theorem 1 (Transaction-cost-safe voluntary exit). *Suppose a position is entered at τ_b and voluntarily exited at a finite time τ_s by either the funnel exit or the trailing-profit exit. If the exit rule includes the condition*

$$X_{\tau_s} - X_{\tau_b} > h,$$

then the after-cost realized multiplier satisfies

$$\frac{(1 - c_{\text{sell}})S_{\tau_s}}{(1 + c_{\text{buy}})S_{\tau_b}} > 1.$$

Proof. By definition of h ,

$$X_{\tau_s} - X_{\tau_b} > \log \left(\frac{1 + c_{\text{buy}}}{1 - c_{\text{sell}}} \right).$$

Exponentiating gives

$$\frac{S_{\tau_s}}{S_{\tau_b}} > \frac{1 + c_{\text{buy}}}{1 - c_{\text{sell}}}.$$

Multiplying by $(1 - c_{\text{sell}})/(1 + c_{\text{buy}})$ proves the claim. $\square$

Remark 2. *The theorem only applies to completed voluntary exits. It does not apply to forced liquidation at a finite horizon, margin liquidation, inability to execute, taxes, bid-ask slippage beyond the modeled cost, or a position that remains open indefinitely.*

6 Wealth and Performance Diagnostics

Let W_t^r denote realized wealth, updated only after completed exits. While holding a position entered at τ_b , the implementation reports after-cost mark-to-market wealth as

$$W_t^m = W_{\tau_b}^r \frac{(1 - c_{\text{sell}})S_t}{(1 + c_{\text{buy}})S_{\tau_b}}.$$

The buy-and-hold benchmark is computed from the same initial exposure convention:

$$W_t^{\text{bh}} = W_0 \frac{(1 - c_{\text{sell}})S_t}{(1 + c_{\text{buy}})S_{\tau_0}}.$$

Let

$$N_{\text{trades}} = \#\{i : \tau_{s,i} \leq T\}$$

be the number of completed trades by horizon T . For completed trades, define the average holding time

$$\bar{\tau}_{\text{hold}} = \frac{1}{N_{\text{trades}}} \sum_{i=1}^{N_{\text{trades}}} (\tau_{s,i} - \tau_{b,i}),$$

when $N_{\text{trades}} > 0$.

The dashboard reports realized return, mark-to-market return, buy-and-hold return, benchmark gap,

$$G_T = \frac{W_T^m}{W_T^{\text{bh}}} - 1,$$

maximum drawdown,

$$D_{\min} = \min_{0 \leq t \leq T} \left(\frac{W_t}{\max_{0 \leq s \leq t} W_s} - 1 \right),$$

exposure, N_{trades} , $\bar{\tau}_{\text{hold}}$, Sharpe ratio, and Calmar ratio. For bar frequency B periods per year, the annualized Sharpe ratio is

$$\text{Sharpe} = \sqrt{B} \frac{\bar{r}}{s_r}, \quad r_t = \frac{W_t - W_{t-1}}{W_{t-1}},$$

and the Calmar ratio is

$$\text{Calmar} = \frac{R_{\text{ann}}}{|D_{\min}|}.$$

The implementation uses $B = 252$ for daily bars, $B = 252 \cdot 6.5$ for hourly bars, and $B = 252 \cdot 26$ for 15-minute bars. These annualization constants assume U.S. equity trading hours; they should be changed for assets or venues with different trading calendars.

7 Adaptive Estimation

The trading rule uses estimates available before observing the next return, preventing lookahead in the displayed funnel.

7.1 Drift

Let $r_t = X_t - X_{t-1}$. The drift is modeled as a slowly moving latent state,

$$\mu_t = \mu_{t-1} + \eta_t, \quad \eta_t \sim \mathcal{N}(0, q),$$

with observation equation

$$r_t = \mu_t + \xi_t, \quad \xi_t \sim \mathcal{N}(0, \sigma^2).$$

The Kalman one-step prediction is

$$\hat{\mu}_{s|s-1} = \hat{\mu}_{s-1|s-1}, \quad P_{s|s-1} = P_{s-1|s-1} + q,$$

and the update after observing r_s is

$$K_s = \frac{P_{s|s-1}}{P_{s|s-1} + \sigma^2},$$

$$\hat{\mu}_{s|s} = \hat{\mu}_{s|s-1} + K_s(r_s - \hat{\mu}_{s|s-1}), \quad P_{s|s} = (1 - K_s)P_{s|s-1}.$$

The UI drift-adaptation control changes q : larger values make the drift estimate react faster to recent returns.

7.2 Volatility

The conditional variance forecast uses a GARCH-style recursion. With innovation

$$e_s = r_s - \hat{\mu}_{s|s-1},$$

define

$$\hat{\sigma}_{s|s-1}^2 = h_s, \quad h_{s+1} = \omega + \alpha e_s^2 + \beta h_s,$$

where $\alpha, \beta \geq 0$ and $\alpha + \beta < 1$. The implementation chooses

$$\omega = (1 - \alpha - \beta)\bar{\sigma}^2$$

from an initial variance estimate $\bar{\sigma}^2$. The funnel variance accumulates the one-step forecasts,

$$V_\tau(t) = \sum_{s=\tau+1}^t \hat{\sigma}_{s|s-1}^2.$$

The Extensions section describes an optional regime-filtering layer based on a three-state hidden Markov model. In that mode, the same trading rule is retained but the one-step funnel inputs are filtered through latent regime probabilities.

8 Backtesting and Real-Data Implementation

The offline walk-forward procedure selects parameters using only past data. For fold j , let $\mathcal{T}_j$ be the training window and $\mathcal{O}_j$ the following out-of-sample window. For candidate parameter vector θ , the training objective used in the implementation has the form

$$J_j(\theta) = R_{j,\text{train}}^m(\theta) + D_{j,\text{train}}^m(\theta) + \lambda G_{j,\text{train}}(\theta),$$

where R^m is mark-to-market return, $D^m \leq 0$ is maximum drawdown, and G is the benchmark gap. The selected parameter is

$$\hat{\theta}_j = \arg \max_{\theta \in \Theta} J_j(\theta).$$

It is then evaluated on $\mathcal{O}_j$ without using future information. The current implementation uses $\lambda = 0.15$, so benchmark-relative performance influences selection but does not dominate absolute mark-to-market return or drawdown.

In the accompanying calibration reports, the parameter search has been explored up to four dimensions: the baseline two-dimensional search over (k, δ) , three-dimensional expansions (k, δ, x) , and four-dimensional expansions (k, δ, x, y) for selected design parameters. Moving beyond four dimensions is possible in principle but becomes computationally expensive and less interpretable. The current project therefore treats four-dimensional calibration as a practical upper bound for Layer 2 sensitivity analysis, rather than as a claim that the full parameter vector has been globally optimized.

For real market data, the app uses a continuous online variant. Instead of resetting wealth and position at every fold, it keeps one account alive. Every update window, recent bars generate a suggested parameter vector $\tilde{\theta}_j$, and the live parameter is smoothed:

$$\theta_j = (1 - \rho)\theta_{j-1} + \rho\tilde{\theta}_j, \quad 0 < \rho \leq 1.$$

Integer parameters such as δ are rounded and constrained to the candidate grid. The current implementation uses learning rate $\rho = 0.25$ and update window 5 bars. This gives a more realistic animation: the account learns from recent data while position state, realized wealth, and mark-to-market wealth continue without reset.

The implementation can fetch Yahoo Finance chart data for several symbols, including SPY, NVDA, TSLA, MSFT, AAPL, AMZN, GOOGL, META, and custom tickers. The selected horizon

determines the bar interval:

lookback	bar interval
5 years	1 day
2 years	1 day
1 year	1 day
1 month	1 hour
1 week	15 minutes

This avoids having only a few observations in short-horizon animations. Intraday results should be interpreted cautiously because microstructure noise, bid-ask spreads, execution latency, and Yahoo interval availability matter more at shorter bar lengths.

9 Empirical Illustration

As a basic Layer 2 validation, the current implementation was run on SPY using Yahoo adjusted price data, a 5-year lookback, 1-day bars, zero transaction costs, drift-adaptation variance $q = 10^{-7}$ (the default UI value), $\lambda = 0.15$, $\rho = 0.25$, $L = 60$, $a = 0.04$, and $z_{\text{trend}} = 0.35$. The purpose of this table is not to claim profitability, but to show that the framework produces the diagnostics required to compare the strategy against buy-and-hold. Because the data window is fetched relative to the current date and the UI sliders remain editable, small differences can occur if the run date or parameters differ.

Quantity	SPY result
Bars	1255
Start price	388.18
End price	720.65
Realized return	53.10%
Mark-to-market return	65.93%
Buy-and-hold return	85.65%
Benchmark gap	−10.62%
Strategy max drawdown	−24.64%
Buy-and-hold max drawdown	−24.50%
Exposure	92.60%
N_{trades}	9
$\bar{\tau}_{\text{hold}}$	117.11 bars
Sharpe ratio	0.69
Calmar ratio	0.43
Online parameter updates	250
Final (k, δ)	(0.80, 2)

In this example, the strategy earned a positive mark-to-market return but underperformed buy-and-hold. This is a useful outcome rather than a failure of the framework: it demonstrates why benchmark gap, drawdown, exposure, trade count, Sharpe ratio, and Calmar ratio must be reported together. The no-realized-loss theorem concerns completed voluntary exits only; it does not imply benchmark outperformance.

10 Extensions

The first implemented Layer 3 extensions are generalized momentum and a three-state regime HMM filter. Both are available in the GUI for simulated paths and real Yahoo Finance data, so they can be compared against the fixed baseline and the online optimization mode on the same selected price path.

10.1 Generalized Momentum

Generalized momentum replaces the fixed exit threshold $M_t \leq 0$ with the holding-time-dependent threshold

$$f(t; \tau_b) = -\frac{c}{\sqrt{t - \tau_b}}, \quad c \geq 0.$$

The baseline rule is recovered when $c = 0$. For $c > 0$, early exits require a stronger negative momentum signal, while later exits gradually relax back toward the original $M_t \leq 0$ condition. The funnel condition $Z_{\tau_b, t} > k$ and the transaction-cost profit buffer $X_t - X_{\tau_b} > h$ are unchanged, so the no-realized-loss theorem still applies to completed voluntary exits. The parameter c is tunable in the GUI.

10.2 Regime HMM Filtering

The HMM extension is best viewed as a filtering layer for the funnel, not as a separate trading rule. The buy, sell, re-entry, and profit-buffer logic remain the same. What changes is the estimate of the next-step drift and variance used inside $C_\tau(t)$ and $V_\tau(t)$.

Let

$$R_t \in \{1, 2, 3\}$$

denote a latent regime, interpreted in the implementation as calm, mixed, and stress states. Conditional on $R_t = j$, returns satisfy

$$r_t \mid R_t = j \sim \mathcal{N}(\mu_j, \sigma_j^2),$$

and R_t evolves according to a persistent Markov transition matrix. The implementation initializes (μ_j, σ_j^2) from an early training segment by sorting returns into volatility buckets, then updates regime probabilities online by forward filtering.

Let

$$\pi_{t|t-1}(j) = P(R_t = j \mid r_{1:t-1})$$

be the predicted regime probability before observing r_t . The filtered one-step drift input is the mixture mean

$$\hat{\mu}_{t|t-1}^{\text{HMM}} = \sum_{j=1}^3 \pi_{t|t-1}(j) \mu_j,$$

with filtered mixture variance

$$\hat{\sigma}_{t|t-1}^{2, \text{HMM}} = \sum_{j=1}^3 \pi_{t|t-1}(j) (\sigma_j^2 + \mu_j^2) - \left(\hat{\mu}_{t|t-1}^{\text{HMM}} \right)^2.$$

These filtered inputs replace the Kalman/GARCH inputs inside the diagnostic funnel. The GUI also displays the most likely regime and posterior probability during playback, making the extension interpretable as the path evolves.

Further Layer 3 work may apply the same time-dependence idea to z_{trend} , making trend re-entry stricter or looser depending on time since the last sale. These additions introduce new calibration degrees of freedom and should be tested by walk-forward validation before being promoted to the main rule.

11 Limitations

The strategy is not risk-free. The main limitations are:

- **Mark-to-market risk:** open positions can lose value before an exit occurs.
- **Holding-time risk:** the algorithm may hold a losing or stagnant position for a long time.
- **Forced-exit risk:** liquidation outside the voluntary exit rule can realize losses.
- **Opportunity cost:** the strategy may underperform buy-and-hold, especially in persistent trends or if re-entry is too conservative.
- **Estimation error:** drift is difficult to estimate, and volatility models are misspecified in real markets.
- **Data and execution risk:** public chart data may contain gaps or limits, and the model does not fully include spreads, taxes, market impact, or order execution.
- **Parameter sensitivity:** k , δ , q , L , a , and z_{trend} affect both returns and risk.

12 Conclusion

The Layer 1–2 strategy now consists of a formal long-only trading rule, an implementation-aligned no-realized-loss theorem for voluntary exits, adaptive drift and volatility estimation, practical bounded and trailing exit logic, two-mode re-entry, walk-forward evaluation, and real-data testing across multiple lookbacks and bar intervals. The guarantee is intentionally limited but useful: completed voluntary exits are transaction-cost-safe by construction. The remaining question is empirical quality, which must be judged through benchmark-relative return, mark-to-market draw-down, exposure, trade frequency, Sharpe ratio, Calmar ratio, and robustness across assets and sampling frequencies.